

Math 56 Compu & Expt Math, Spring 2013: HW6 Debriefing

1. $2+3 = 5$ pts

- (a) e^{-n^γ} for any $0 < \gamma < 1/2$ works, but Kyutae invented a different superalgebraic function $n^{\log n}$. A complete answer proves your func is superalgebraic yet slower converging than $e^{-c\sqrt{n}}$ for any $c > 0$.
- (b) You all found the max n occurs at $n^2 = 2k/c$. Treating n as a real number here is a neat trick to bound its value on the integers. The point here is that superalgebraic doesn't mean the constant C_k has to be bounded uniformly in the order k .

2. $5+4 = 9$ pts

- (a) You just computed a number with over one million digits in a second or two! The joy of FFTs. A crucial step is to **round** the real-valued result from the FFT to integers before the carrying is done, otherwise when the test of whether to carry is done, eg 9.9999999 is treated as 9 not as 10.
- (b) For the time estimate of the naive method (mult by 2, repeat 2^{22} times), each mult is $O(D)$ not $O(D \log D)$ since it's mult by a small $O(1)$ number, where D = number of digits in answer. Overall is $O(D^2) = 2^{44} \approx 10^{13}$ flops, several hours of CPU time. Tom has good analysis.

BONUS See John for proof that the last digit is always 6. See Hanh for last 2 digits on a period-4 orbit. The last 3 digits happen to form a period-20 orbit (I'm not sure of significance of this number—are you?)

3. $4+2+2+3 = 11$ pts. Some of the early points for getting into python.

- (a) As several found, if the answer is x , initial guesses must be in $(0, 2x)$ otherwise the iteration blows up to infinity. See eg Kyutae discussion.
- (b) The repeating string is 96 long:
010309278350515463917525773195876288659793814432989690721649484536082474226804123711340206185567
Hanh explains this in terms of Fermat's Little Theorem. Also see
http://en.wikipedia.org/wiki/Repeating_decimal
- (c) $O(N \log^2 N)$, see eg John.
- (d) As you found, the errors hardly differ, and in fact the first way is slightly *more* accurate. Eg see Kyutae. Wikipedia is wrong! (Please, someone correct it)

4. $3+4+2+4 = 13$ pts.

Throughout this question it was important to *check convergence* to the required accuracy! One way to guarantee this is a **while** loop that only stops when the answer doesn't change to the required precision. Another (harder) is to precompute the number of terms using a convergence rate or estimate.

- (a) The definition of the number of “terms used” is a bit ambiguous, since it could mean “what is the highest power n of x used in the Taylor series?” By that definition, as Kunyi calculates, if $\tan^{-1} 1/5$ is the largest number to approximate by a Taylor series, around $n = 14500$ is needed. But this involves only $n/2$ nonzero terms, i.e. around 7300 “terms.” Either definition I treated as correct.

- (b) As many of you discovered, online tools allow you to check digits of π , or, better, sage or python/mpmath can do it efficiently via `mp.dps = 10000; print str(pi)[-10:]` which prints digits 9991 to 10000. Note the python array indexing (equivalent to `(end-9:end)` in Matlab). We believe mpmath uses Brent–Salamin. If you rounded the 10000th digit from 7 to 8, this was fine. Careful with defining modules with names like `pi`, etc. This overwrites the constant `pi` that mpmath has!
- (c) Difference of two squares; review of quadratic convergence proofs. Hanh shows that since the iterations x_n, y_n always lie between the original values, you can use $\min(x_0, y_0)$ to bound the const. Proving the convergence of α_n would be extra; see Salamin’s original 1976 paper.
- (d) You all found Brent–Salamin around 10^2 times faster than Taylor series even at $N = 10^4$ digits. Imagine how much faster it is at $N = 10^6$. Your algorithms were close to mpmath’s in speed for evaluating π , ie one million digits in a couple of seconds! See eg Kunyi for the digits.